

Pandas

March 1, 2024

- index -> axis=0
- columns -> axis=1

1 Initialisation

```
[1]: import pandas as pd
import numpy as np
from IPython.display import HTML, Markdown, Latex

def display_df(tp_df=None, index=False):
    # tp_df = tp_df if isinstance(tp_df, pd.DataFrame) else df
    tp_df = tp_df if tp_df is not None else df
    display(Markdown(tp_df.to_markdown(index=index)))

# def display_df(tp_df=None, index=False):
#     tp_df = tp_df if isinstance(tp_df, pd.DataFrame) else df
#     display(HTML(tp_df.to_html(index=index)))
```

2 Creating a dataframe

```
[2]: data = {
    'age': [10, 22, 13, 21, 12, 11, 17],
    'section': ['A', 'B', 'C', 'B', 'B', 'A', 'A'],
    'city': ['Gurgaon', 'Delhi', 'Mumbai', 'Delhi', 'Mumbai', 'Delhi', 'Mumbai'],
    'gender': ['M', 'F', 'F', 'M', 'M', 'M', 'F'],
    'favorite_color': ['red', 'black', 'yellow', 'pink', 'black', 'green', 'red']
}

data_csv = [
    ['age', 'section', 'city', 'gender', 'favorite_color'],
    [10, 'A', 'Gurgaon', 'M', 'red'],
    [22, 'B', 'Delhi', 'F', 'black'],
    [13, 'C', 'Mumbai', 'F', 'yellow'],
    [21, 'B', 'Delhi', 'M', 'pink'],
    [12, 'B', 'Mumbai', 'M', 'black']
]
```

```

    [11,      'A',      'Delhi',   'M',      'green'      ],
    [17,      'A',      'Mumbai',  'F',      'red'        ]
]

data_dict = [
    {'age': 10, 'section': 'A', 'city': 'Gurgaon', 'gender': 'M',
    ↪ 'favorite_color': 'red'},
    {'age': 22, 'section': 'B', 'city': 'Delhi',   'gender': 'F',
    ↪ 'favorite_color': 'black'},
    {'age': 13, 'section': 'C', 'city': 'Mumbai',  'gender': 'F',
    ↪ 'favorite_color': 'yellow'},
    {'age': 21, 'section': 'B', 'city': 'Delhi',   'gender': 'M',
    ↪ 'favorite_color': 'pink'},
    {'age': 12, 'section': 'B', 'city': 'Mumbai',  'gender': 'M',
    ↪ 'favorite_color': 'black'},
    {'age': 11, 'section': 'A', 'city': 'Delhi',   'gender': 'M',
    ↪ 'favorite_color': 'green'},
    {'age': 17, 'section': 'A', 'city': 'Mumbai',  'gender': 'F',
    ↪ 'favorite_color': 'red'}
]

df = pd.DataFrame(data)
display_df()

df = pd.DataFrame(data_csv[1:], columns=data_csv[0])
display_df()

df = pd.DataFrame(data_dict)
display_df()

```

age	section	city	gender	favorite_color
10	A	Gurgaon	M	red
22	B	Delhi	F	black
13	C	Mumbai	F	yellow
21	B	Delhi	M	pink
12	B	Mumbai	M	black
11	A	Delhi	M	green
17	A	Mumbai	F	red

age	section	city	gender	favorite_color
10	A	Gurgaon	M	red
22	B	Delhi	F	black
13	C	Mumbai	F	yellow
21	B	Delhi	M	pink

age	section	city	gender	favorite_color
12	B	Mumbai	M	black
11	A	Delhi	M	green
17	A	Mumbai	F	red

age	section	city	gender	favorite_color
10	A	Gurgaon	M	red
22	B	Delhi	F	black
13	C	Mumbai	F	yellow
21	B	Delhi	M	pink
12	B	Mumbai	M	black
11	A	Delhi	M	green
17	A	Mumbai	F	red

3 Meta stuff

```
[3]: data = {
    'age':          [10,22,13,21,12,11,17],
    'section':      ['A','B','C','B','B','A','A'],
    'city':         ['Gurgaon','Delhi','Mumbai','Delhi','Mumbai','Delhi','Mumbai'],
    'gender':       ['M','F','F','M','M','M','F'],
    'favorite_color': ['red','black','yellow','pink','black','green','red']
}
```

```
df = pd.DataFrame(df)
```

```
print(f'df.empty:    {df.empty}')
print(f'df.shape:    {df.shape}')
print(f'df.index:    {df.index}')
print(f'df.columns:  {df.columns}')
```

```
print('\ndf.describe():')
display_df(df.describe(), index=True)
```

```
print('\ndf.info():')
df.info() # This automatically prints stuff to stdout
```

```
df.empty:    False
df.shape:    (7, 5)
df.index:    RangeIndex(start=0, stop=7, step=1)
df.columns:  Index(['age', 'section', 'city', 'gender', 'favorite_color'],
dtype='object')
```

```
df.describe():
```

	age
count	7
mean	15.1429
std	4.8795
min	10
25%	11.5
50%	13
75%	19
max	22

```
df.info():
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7 entries, 0 to 6
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   age              7 non-null      int64
1   section          7 non-null      object
2   city             7 non-null      object
3   gender           7 non-null      object
4   favorite_color   7 non-null      object
dtypes: int64(1), object(4)
memory usage: 412.0+ bytes
```

4 .head() and .tail()

Both take a single optional argument, ie `n`, which is an integer representing the number of records to show. By default, it is 5. `.head()` shows the `n` top most records and `.tail()` shows the `n` bottom most records

```
[4]: display_df( df.head(3) , index=True )
      display_df( df.tail(3) , index=True )
```

	age	section	city	gender	favorite_color
0	10	A	Gurgaon	M	red
1	22	B	Delhi	F	black
2	13	C	Mumbai	F	yellow

	age	section	city	gender	favorite_color
4	12	B	Mumbai	M	black
5	11	A	Delhi	M	green
6	17	A	Mumbai	F	red

5 .iloc[]

Format is something like `DataFrame.iloc[row_indexer, column_indexer]`. Here `column_indexer` is optional. If it is not given, all columns will be printed

Remember, here rows start with index 1 rather than index 0

- `row_indexer`: This can be a slice (like `1:9:2` or `:`), or a list of the indexes, like `[1,4,5]`
- `column_indexer`: This can be a slice (like `1:9:2` or `:`), or a list of the indexes of the columns, like `[1,4,5]`

```
[5]: data = {
      'age':          [10,22,13,21,12,11,17],
      'section':      ['A','B','C','B','B','A','A'],
      'city':         ['Gurgaon','Delhi','Mumbai','Delhi','Mumbai','Delhi','Mumbai'],
      'gender':       ['M','F','F','M','M','M','F'],
      'favorite_color': ['red','black','yellow','pink','black','green','red']
    }
df = pd.DataFrame(data)

print('\nOriginal data:')
display_df()

print('\nRecords of index 1 & 3')
display_df( df.iloc[ [1,3] , : ] )
```

Original data:

	age	section	city	gender	favorite_color
	10	A	Gurgaon	M	red
	22	B	Delhi	F	black
	13	C	Mumbai	F	yellow
	21	B	Delhi	M	pink
	12	B	Mumbai	M	black
	11	A	Delhi	M	green
	17	A	Mumbai	F	red

Records of index 1 & 3

	age	section	city	gender	favorite_color
	22	B	Delhi	F	black
	21	B	Delhi	M	pink

6 .loc[]

Just like `.iloc[]`, but uses names (strings) rather than indexes, and the slicing is end-inclusive (unlike the slicing we have seen till now)

```
[ ]:
```

7 .rename()

It by default doesn't do the renaming inplace, and instead returns a copy

(Here the `mapper`, `index`, and `columns` arguments accept a dict-like object or a function)

Arguments:

- `mapper`: Dict-like or function transformations to apply to that axis' values
- `axis`: `int | str = 0`: Axis to target with mapper. Can be 0/`index` or 1/`columns`
- `index`: Alternative to specifying axis (`mapper,axis=0` is equivalent to `index=mapper`)
- `columns`: Alternative to specifying axis (`mapper,axis=1` is equivalent to `columns=mapper`)
- `inplace`: `bool = False`: Whether to modify the DataFrame rather than creating a new one
- `errors`: `str`: Can be "raise" or "ignore". If "raise", then raise a `KeyError` when a dict-like mapper, index, or columns contains labels that are not present in the Index being transformed. If "ignore", existing keys will be renamed and extra keys will be ignored

```
[6]: df = pd.DataFrame({'A': [1, 2, 3], 'B': [4, 5, 6]})

df.rename(columns={'A': 'New A', 'B': 'New B'}, inplace=True)
df.rename(index=lambda x: f'Row {x}', inplace=True)
display_df(index=True)
```

	New A	New B
Row 0	1	4
Row 1	2	5
Row 2	3	6

8 .query()

Can do SQL-python like queries. Returns a dataframe

The format is `DataFrame.query(expr:str, inplace:bool)`. If `inplace` is `True`, `None` is returned and the original df is replaced by the df which would had been returned if `inplace` was `False`

```
[7]: data = {
    'age':          [10,22,13,21,12,11,17],
    'section':      ['A','B','C','B','B','A','A'],
    'city':         ['Gurgaon','Delhi','Mumbai','Delhi','Mumbai','Delhi','Mumbai'],
    'gender':       ['M','F','F','M','M','M','F'],
```

```

    'favorite_color': ['red','black','yellow','pink','black','green','red']
}
df = pd.DataFrame(data)

print('\nOriginal data:')
display_df()

print('\nRecords where age >= 15:')
display_df( df.query('age >= 15') )

print('\nRecords where age >= 12 and gender = Male:')
display_df( df.query('age >= 12 and gender == "M"') )

print('\nCity and gender of people with age >= 12:')
display_df( df.query('age >= 12')[['city','gender']] )

# Use of `@` and ``

```

Original data:

age	section	city	gender	favorite_color
10	A	Gurgaon	M	red
22	B	Delhi	F	black
13	C	Mumbai	F	yellow
21	B	Delhi	M	pink
12	B	Mumbai	M	black
11	A	Delhi	M	green
17	A	Mumbai	F	red

Records where age >= 15:

age	section	city	gender	favorite_color
22	B	Delhi	F	black
21	B	Delhi	M	pink
17	A	Mumbai	F	red

Records where age >= 12 and gender = Male:

age	section	city	gender	favorite_color
21	B	Delhi	M	pink
12	B	Mumbai	M	black

City and gender of people with age >= 12:

city	gender
Delhi	F
Mumbai	F
Delhi	M
Mumbai	M
Mumbai	F

9 .sort_values()

Arguments:

- **by:** str | list[str]: Name or list of names to sort by
- **axis:** int | str = 0: Axis to be sorted. Can be 0/index or 1/columns
- **ascending:** bool | list[bool] = True: Self explanatory. Specify list for multiple sort orders. If this is a list of bools, must match the length of the by
- **inplace:** bool = False: If True, perform operation in-place
- **na_position:** str = "last": Puts NaNs at the beginning if **first**, and at the end if **last**
- **ignore_index:** bool = False: If True, the resulting axis will be labeled 0,1,...,n-1

```
[8]: data = {
    'age':          [10,22,13,21,12,11,17],
    'section':      ['A','B','C','B','B','A','A'],
    'city':         ['Gurgaon','Delhi','Mumbai','Delhi','Mumbai','Delhi','Mumbai'],
    'gender':       ['M','F','F','M','M','M','F'],
    'favorite_color': ['red','black','yellow','pink','black','green','red']
}
df = pd.DataFrame(data)

print('\nOriginal data:')
display_df()

print('\nSorted by age (descending):')
display_df( df.sort_values(by='age',ascending=False).head(3) , index=True )
```

Original data:

age	section	city	gender	favorite_color
10	A	Gurgaon	M	red
22	B	Delhi	F	black
13	C	Mumbai	F	yellow
21	B	Delhi	M	pink

age	section	city	gender	favorite_color
12	B	Mumbai	M	black
11	A	Delhi	M	green
17	A	Mumbai	F	red

Sorted by age (descending):

	age	section	city	gender	favorite_color
1	22	B	Delhi	F	black
3	21	B	Delhi	M	pink
6	17	A	Mumbai	F	red

10 `.count()`, `.sum()`, `min()`, `max()`, `.mean()`, `median()`, and `mode()`

All of these return a `pd.core.series.Series`, except `.mode()`. Mode returns a `df` cause there might be many values that are the mode, and different rows contains these different values

10.0.1 `.count()`

- `axis: int | str = 0`: If 0 or “index” counts are generated for each column. If 1 or “columns”, counts are generated for each row
- `numeric_only: bool = False`: Include only float, int or boolean data‘

10.0.2 `.sum()`

- `axis: int | str = 0`: Axis for the function to be applied on
- `numeric_only: bool = False`: Include only float, int or boolean data
- `skipna: bool = True`: Exclude NA/null values when computing the result
- `min_count: int = 0`: The required number of valid values to perform the operation. If fewer non-NA values are present, the result will be NA

10.0.3 `.min()`, `.max()`, `.mean()`, `.median()`, & `.mode()`

- `axis: int | str = 0`: Axis for the function to be applied on
- `numeric_only: bool = False`: Include only float, int or boolean data
- `skipna: bool = True`: Exclude NA/null values when computing the result

```
[9]: data = {
      "Person": ["John", "Myla", "Lewis", "John", "Myla"],
      "Age":    [24, np.nan, 21, 33, 26],
      "Single": [False, True, True, True, False]
    }
    df = pd.DataFrame(data)

    print('Original df:')
    display_df()
```

```

print('\n1) Count:')
display_df( df.count(), index=True)
display_df( df.count(numeric_only=True), index=True)

print('\n2) Sum:')
display_df( df.sum(), index=True)

print('\n3) Min:')
display_df( df.min(), index=True)

print('\n4) Max:')
display_df( df.max(), index=True)

print('\n5) Mean:')
display_df( df.mean(numeric_only=True), index=True)

print('\n6) Median:')
display_df( df.median(numeric_only=True), index=True)

print('\n7) Mode:')
display_df( df.mode(numeric_only=True), index=True)

```

Original df:

Person	Age	Single
John	24	False
Myla	nan	True
Lewis	21	True
John	33	True
Myla	26	False

1) Count:

	0
Person	5
Age	4
Single	5
	0
Age	4
Single	5

2) Sum:

	0
Person	JohnMylaLewisJohnMyla
Age	104.0
Single	3

3) Min:

	0
Person	John
Age	21.0
Single	False

4) Max:

	0
Person	Myla
Age	33.0
Single	True

5) Mean:

	0
Age	26
Single	0.6

6) Median:

	0
Age	25
Single	1

7) Mode:

	Age	Single
0	21	1
1	24	nan
2	26	nan

	Age	Single
3	33	nan

- 11 `.groupby()`
- 12 `.select_dtypes()`
- 13 `.duplicated()` and `drop_duplicates()`
- 14 `pd.concat()` and `pd.append()`
- 15 `.pivot()` and `pivot_table()`